

Exploratory Data Analysis: NZ Crash Data

Introduction

This is an Exploratory Data Analysis (EDA) for New Zealand Crash data. The aim of EDA is to explore data, find interesting patterns, and generate set of hypothesis that are worth to explore further, but also to find inconsistencies, coding errors, duplications and so on. For this reason, EDA is often the first step that needs to be done when looking at new data before other statistical or machine learning approaches are used. Ideally, EDA will be done in collaboration with a domain expert as many of the patterns or inconsistencies can be just a property of the data, be it from the way the data was collected, or expected patterns based on the mechanism that generate the data. Since these kind of patterns might influence the model-fitting and hypothesis testing steps later in the analysis, identifying these patterns and working with domain expert is an important step to get aware of these patterns so that they do not influence steps later and do not produce false positive relationships, or mask the true ones.

EDA is not a hard science, but more of an art, and there are no hard rules. While some tools for automatic EDA exist, see *The Landscape of R Packages for Automated Exploratory Data Analysis* for review, EDA is in its essence interactive. You are going through data, looking at various data points, features, and patterns, and searching for interesting questions and answers.

The history of EDA is interesting and its modern history starts with Tukey, but in truth, any data analysis starts with looking at patterns in the data and in truth, EDA is as old as statistics, and perhaps even predates it. After all, this is how Bayes came with, and Laplace rediscovered the Bayesian theorem, or how various probability distributions were invented/discovered. For a good review on statistics data science, and EDA, see *50 years of Data Science*. Or you can go directly for the *Future of Data Analysis* from Tukey written in 1961. For a good tutorial for EDA, the *R for Data Science* from Hadley is a great source.

EDA: NZ Crash Data

The purpose of this analysis is to explore New Zealand crash statistics data, and create insight on crashes by region, or other spatial area. And if we find some interesting relationship, we might try to do some little non-serious modelling. Why non-serious? Because we need domain expert to discuss the data with us, but we do not have a domain expert. Because of that, we do not know what are the limitations and assumptions in the data, so any modelling done in here is purely exploratory.

As pointed by the previous paragraph, there are some limitations. The challenge here is to do this in a limited time, so we will be able to do only very cursory exploration and look only at some overall patterns in the data, we won't be going much into detail both by looking at year to year differences in various statistics, or by looking closely at small geographical areas, such as spikes in the city centres and so on. We also do not have a domain expert on hand, so we can't discuss various findings, which means we do not know if any pattern we see is a true pattern, or just something resulting from the way the data was collected/encoded. This will be obvious quite later.

Finally, as this is exploratory analysis, do not expect high quality infographics. Quick and dirty graphs are what we are looking for. In fact, in some some cases, I prefer to not even look at graphs, instead look at distribution of numbers. Good graphs take a time to setup, and sometimes might hide some obvious issues with the data. For instance, boxplots are great for visualising basic statistics, but it hides how the data are distributed inside the box, for instance, perfectly clean boxplot might not show a bimodal distribution of data.

Downloading the data

To get the data, we download the CSV from <https://opendata-nzta.opendata.arcgis.com>.

I downloaded a table updated on 6th July 2023, and there are 821 744 records. Unfortunately, there is no way how to provide a direct link to download the data, as this is hidden behind javascript action. There is a ArcGIS REST API, but it is not documented at this page, and even the main documentation, once you figure out that you are looking for **Feature Service**, is less than ideal.

Given enough time, this could be figured out, but by default, only 2000 records are downloaded, and since we are pressed for time, I chose to not explore this further. Last time I worked with a similar API (GenBank), there was a hard limit on how many records could be downloaded at once. To download everything, one had to split the query into multiple subqueries and collate the result. To do this right, I would need to either use or write a package to provide binding for these queries and this would take quite a bit of time. It is possible to do this without package, but it might get quite a bit difficult and error-prone.

Basic data cleaning and exploration

I downloaded the data into the `data` folder, so we can start looking at it and doing some basic cleaning.

```
filename = "data/Crash_Analysis_System_(CAS)_data.csv"
data = read.csv(filename)
```

```
names(data)
```

```
## [1] "X"
## [3] "OBJECTID"
## [5] "areaUnitID"
## [7] "bridge"
## [9] "carStationWagon"
## [11] "crashDirectionDescription"
## [13] "crashLocation1"
## [15] "crashRoadSideRoad"
## [17] "crashSHDescription"
## [19] "debris"
## [21] "ditch"
## [23] "fence"
## [25] "guardRail"
## [27] "houseOrBuilding"
## [29] "kerb"
## [31] "meshblockId"
## [33] "moped"
## [35] "NumberOfLanes"
## [37] "otherObject"
## [39] "overBank"
## [41] "pedestrian"
## [43] "postOrPole"
## [45] "roadCharacter"
## [47] "roadSurface"
## [49] "schoolBus"
## [51] "slipOrFlood"
## [53] "strayAnimal"
## [55] "suv"
## [57] "temporarySpeedLimit"
## [59] "tlaName"
## [61] "trafficIsland"
## [63] "train"
## [1] "Y"
## [3] "advisorySpeed"
## [5] "bicycle"
## [7] "bus"
## [9] "cliffBank"
## [11] "crashFinancialYear"
## [13] "crashLocation2"
## [15] "crashSeverity"
## [17] "crashYear"
## [19] "directionRoleDescription"
## [21] "fatalCount"
## [23] "flatHill"
## [25] "holiday"
## [27] "intersection"
## [29] "light"
## [31] "minorInjuryCount"
## [33] "motorcycle"
## [35] "objectThrownOrDropped"
## [37] "otherVehicleType"
## [39] "parkedVehicle"
## [41] "phoneBoxEtc"
## [43] "region"
## [45] "roadLane"
## [47] "roadworks"
## [49] "seriousInjuryCount"
## [51] "speedLimit"
## [53] "streetLight"
## [55] "taxi"
## [57] "tlaId"
## [59] "trafficControl"
## [61] "trafficSign"
## [63] "tree"
```

```
## [65] "truck"                "unknownVehicleType"
## [67] "urban"                 "vanOrUtility"
## [69] "vehicle"               "waterRiver"
## [71] "weatherA"              "weatherB"
```

There is quite lot of features in here. One big issue with this dataset is that there is a complete lack of description *what these features mean*. Previously, I worked with Ethnographic Atlas, and there is a whole codebook on the meaning of individual features and the way they were coded. IMHO, this is quite an issue that should be corrected and something that could be easily resolved if we had an expert on hand.

For instance, what is `advisorySpeed`? Advised speed lower than maximum allowed speed that is put in tight corners? That would make sense. But look at what happens when we create a contingency table (bivariate relationship) of `advisorySpeed` and `speedLimit`, which I assume is the maximum allowed speed:

```
table("Speed Limit"=data$SpeedLimit, "Advisory Speed"=data$advisorySpeed, useNA="always")
```

```
##           Advisory Speed
## Speed Limit  15    20    25    30    35    40    45    50    55
##      2      0     0     0     0     0     0     0     0     0
##      5      0     0     0     0     0     0     0     0     0
##      6      0     0     0     0     0     0     0     0     0
##     10      0     0     0     0     0     0     0     0     0
##     15      0     0     0     0     0     0     0     0     0
##     20     11     1     0     0     0     0     0     0     0
##     30     24    20    33     0     0     0     0     0     0
##     40      2     2     5     3     6     0     0     1     0
##     50    172   353   637    915   1220   498   1254     7     2
##     51      0     0     0     0     0     0     0     0     0
##     60      6     1    25    11    99    11    120    27   119
##     61      0     0     0     0     0     0     0     0     0
##     70      4     7    70    29   224    20    201   102   207
##     80      6     8   101    35   444     8    463    83   598
##     90      0     0     0     0     0     0     3     0     1
##    100     65    27   658    315  2053    88   2802   421  3750
##    110      0     0     0     0     2     0     1     0     0
##    <NA>      0     0     0     0     0     0     0     0     0
##           Advisory Speed
## Speed Limit  60    65    70    75    80    85    90    95    <NA>
##      2      0     0     0     0     0     0     0     0     1
##      5      0     0     0     0     0     0     0     0    15
##      6      0     0     0     0     0     0     0     0     1
##     10      0     0     0     0     0     0     0     0   813
##     15      0     0     0     0     0     0     0     0    10
##     20      0     0     0     0     0     0     0     0  1934
##     30      0     0     0     0     0     0     0     0  7261
##     40      0     0     0     0     0     0     0     0  1800
##     50      0     1     0     2     0     0     0     0 488142
##     51      0     0     0     0     0     0     0     0     1
##     60      1     3     0     1     0     0     0     0 22471
##     61      0     0     0     0     0     0     0     0     1
##     70     29    200     1     1     0     1     0     0 24372
##     80     41    488    60   255     0     0     0     0 37847
##     90      0     0     0     4     1     3     0     0   465
##    100    213  4556   501  3684   675  1798   153   289 204254
##    110      0     1     0     0     0     0     0     1   174
##    <NA>      0     0     0     0     0     0     0     0   838
```

Advisory Speed is usually below the Speed Limit, as expected, but I would expect that Advisory Speed is *always* below the Speed Limit. In some cases, this is not true. Are the assumptions we made about the meaning of both variables correct? If so, does that mean that there is a coding error? Or something else is happening, such as incorrect signage on road?

For instance, we could easily subset some of the strange datapoints, for instance, where Speed Limit is 50, and Advisory Speed is larger than that:

```
subset(data, data$speedLimit == 50 & data$advisorySpeed > 50)
```

##		X	Y	OBJECTID	advisorySpeed	areaUnitID	bicycle	bridge	bus
##	98658	1823770	5537280	160173	75	561811	0	0	0
##	192652	1819785	5576514	316844	65	558700	0	NA	0
##	295973	1572148	5170547	480920	55	596400	0	0	0
##	517292	1751258	5950941	822065	55	505805	0	NA	0
##	721471	1571780	5170316	1158036	75	596503	0	0	0
##		carStationWagon	cliffBank	crashDirectionDescription	crashFinancialYear				
##	98658		1	0	South			2020/2021	
##	192652		1	NA	West			2018/2019	
##	295973		0	0	North			2018/2019	
##	517292		2	NA	South			2021/2022	
##	721471		1	1	East			2020/2021	
##		crashLocation1		crashLocation2	crashRoadSideRoad				
##	98658		RAILWAY ROAD	MAPLE STREET	NA				
##	192652		01N-0885	KOTUKUTUKU ROAD	NA				
##	295973		GOVERNORS BAY ROAD	SANDY BEACH ROAD	NA				
##	517292		HIBISCUS COAST HIGHWAY	GRUTS BR	NA				
##	721471		DYERS PASS ROAD	GOVERNORS BAY ROAD	NA				
##		crashSeverity	crashSHDescription	crashYear	debris				
##	98658		Non-Injury Crash	No	2021	0			
##	192652		Minor Crash	Yes	2018	NA			
##	295973		Non-Injury Crash	No	2019	0			
##	517292		Non-Injury Crash	No	2021	NA			
##	721471		Non-Injury Crash	No	2020	0			
##		directionRoleDescription	ditch	fatalCount	fence	flatHill	guardRail		
##	98658		South	0	0	0	Flat	0	
##	192652		South	NA	0	NA	Flat	NA	
##	295973		South	0	0	0	Hill Road	0	
##	517292		South	NA	0	NA	Hill Road	NA	
##	721471		East	0	0	0	Hill Road	0	
##		holiday	houseOrBuilding	intersection	kerb	light	meshblockId		
##	98658		0	NA	0	Dark	1792100		
##	192652		NA	NA	NA	Dark	1728600		
##	295973		0	NA	0	Dark	2711400		
##	517292		NA	NA	NA	Bright sun	171102		
##	721471		0	NA	0	Dark	2711501		
##		minorInjuryCount	moped	motorcycle	NumberOfLanes	objectThrownOrDropped			
##	98658		0	0	0	2		0	
##	192652		1	0	0	2		NA	
##	295973		0	0	0	2		0	
##	517292		0	0	0	2		NA	
##	721471		0	0	0	2		0	
##		otherObject	otherVehicleType	overBank	parkedVehicle	pedestrian			
##	98658		0	0	0	0	NA		
##	192652		NA	0	NA	NA	NA		

```

## 295973      0      0      0      0      NA
## 517292     NA      0     NA     NA     NA
## 721471      0      0      0      0     NA
##      phoneBoxEtc postOrPole      region roadCharacter roadLane
## 98658      0      0 Manawatū-Whanganui Region      Nil 2-way
## 192652     NA     NA Manawatū-Whanganui Region      Nil 2-way
## 295973      0      0      Canterbury Region      Nil 2-way
## 517292     NA     NA      Auckland Region      Nil 2-way
## 721471      0      0      Canterbury Region      Nil 2-way
##      roadSurface roadworks schoolBus seriousInjuryCount slipOrFlood
## 98658     Sealed      0      0      0      0
## 192652     Sealed     NA      0      0     NA
## 295973     Sealed      0      0      0      0
## 517292     Sealed     NA      0      0     NA
## 721471     Sealed      0      0      0      0
##      speedLimit strayAnimal streetLight suv taxi temporarySpeedLimit tlaId
## 98658      50      0      None 0 0      NA 40
## 192652      50     NA      None 0 0      NA 38
## 295973      50      0      On 0 0      NA 60
## 517292      50     NA      None 0 0      NA 76
## 721471      50      0      None 0 0      NA 60
##      tlaName trafficControl trafficIsland trafficSign train
## 98658 Palmerston North City      Nil      0      0 0
## 192652 Rangitikei District      Unknown     NA     NA NA
## 295973 Christchurch City      Unknown      0      0 0
## 517292 Auckland      Give way     NA     NA NA
## 721471 Christchurch City      Nil      0      0 0
##      tree truck unknownVehicleType urban vanOrUtility vehicle waterRiver
## 98658      0 0      0 Urban      0 0 0
## 192652     NA 0      0 Urban      0 NA NA
## 295973      1 0      0 Urban      1 0 0
## 517292     NA 0      0 Urban      0 NA NA
## 721471      0 0      0 Urban      0 0 0
##      weatherA      weatherB
## 98658      Fine      Null
## 192652     Fine      Null
## 295973     Fine      Null
## 517292     Fine Strong wind
## 721471     Fine      Null

```

but I don't see anything particularly wrong or strange.

Knowing this, we will start cleaning the data, and by that I mostly mean dropping variables.

For instance, we won't be working with many geographical variables, and we do not need their IDs, so we can drop `X`, `Y`, `OBJECTID`, `areaUnitID`, `meshblockId`, `crashFinancialYear`, and `tlaId`.

The `tla` stands for Territorial Local Authority, which divides regions into subregions (contained in `region` variable), so we will rename `tlaName` into `subregion` for better understandability.

```

data[c("X", "Y", "OBJECTID", "areaUnitID", "meshblockId", "crashFinancialYear", "tlaId")] = NULL
names(data)[names(data) == "tlaName"] = "subregion"

```

There is more cleaning that we will do, but for that, we need to look at the data. A simple `lapply(data, table)` will do most of the work, since all the data is categorical or ordinal. But doing so for 65 variables is quite a bit too much, so we will divide the features according to their type. We will start with `vehicleType`

and continue with other groups such as `roadConditions`, `otherObjects` and so on.

Also, to help us in this, I define the operator `%-%` so we can see what we already defined.

```
"%-%" = function(x,y){x[!(x %in% y)]}

vehicleType = c("bicycle", "bus", "carStationWagon", "moped", "motorcycle",
  "otherVehicleType", "parkedVehicle", "schoolBus", "suv", "taxi", "train",
  "truck", "unknownVehicleType", "vanOrUtility", "vehicle")

names(data) %-% vehicleType
```

```
## [1] "advisorySpeed"      "bridge"
## [3] "cliffBank"         "crashDirectionDescription"
## [5] "crashLocation1"    "crashLocation2"
## [7] "crashRoadSideRoad" "crashSeverity"
## [9] "crashSHDescription" "crashYear"
## [11] "debris"            "directionRoleDescription"
## [13] "ditch"             "fatalCount"
## [15] "fence"             "flatHill"
## [17] "guardRail"         "holiday"
## [19] "houseOrBuilding"   "intersection"
## [21] "kerb"              "light"
## [23] "minorInjuryCount"  "NumberOfLanes"
## [25] "objectThrownOrDropped" "otherObject"
## [27] "overBank"          "pedestrian"
## [29] "phoneBoxEtc"       "postOrPole"
## [31] "region"            "roadCharacter"
## [33] "roadLane"          "roadSurface"
## [35] "roadworks"         "seriousInjuryCount"
## [37] "slipOrFlood"       "speedLimit"
## [39] "strayAnimal"       "streetLight"
## [41] "temporarySpeedLimit" "subregion"
## [43] "trafficControl"    "trafficIsland"
## [45] "trafficSign"       "tree"
## [47] "urban"             "waterRiver"
## [49] "weatherA"          "weatherB"
```

This helps us filtering down the list of names until we ended with everything categorized into somewhat related variables. It is just for a better systematic exploration, it doesn't have to be perfect. For that, we would need an expert.

```
roadConditions = c("advisorySpeed", "temporarySpeedLimit", "speedLimit", "roadCharacter",
  "roadLane", "roadSurface", "roadworks", "trafficControl", "trafficSign",
  "streetLight", "NumberOfLanes", "intersection", "trafficIsland")

terrainFeature = c("bridge", "cliffBank", "ditch", "fence", "flatHill",
  "guardRail", "houseOrBuilding", "kerb", "overBank",
  "waterRiver", "slipOrFlood")

otherObjects = c("debris", "objectThrownOrDropped", "otherObject", "pedestrian",
  "phoneBoxEtc", "postOrPole", "strayAnimal", "tree")

crashSeverity = c("crashSeverity", "fatalCount", "minorInjuryCount", "seriousInjuryCount")

weather = c("light", "weatherA", "weatherB")
```

```
location = c("crashDirectionDescription", "crashLocation1", "crashLocation2", "directionRoleDescription")

time = c("crashYear", "holiday")

other = c("crashRoadSideRoad", "crashSHDescription")
```

Now, we can look closely at all these variables and their distributions.

Other variables Let's start with the `other` category:

```
# this is very efficient for ordinal or categorical variables:
data[other] |> lapply(table)
```

```
## $crashRoadSideRoad
## < table of extent 0 >
##
## $crashSHDescription
##
##      No Unknown      Yes
## 580900      55 240789
```

The `crashRoadSideRoad` is empty and we can safely drop it.

I didn't know what `crashSHDescription` means, so I had to google around and apparently, there is an old key from 2019! The fields and values differ, but there is at least some explanation. Apparently, this variable indicate if the crash happened on State Highway (SH) or somewhere else.

The encoding of `crashSHDescription` is nice and readable, but if we meant to fit it into model, we would like binary variable with unknown coded as NA instead, so we recode it.

```
data$crashRoadSideRoad = NULL

# probably my favourite way how to do recoding
data$crashSHDescription = setNames(
  c(0,1,NA), c("No", "Yes", "Unknown")
)[data$crashSHDescription]
```

Time variables Now we look at the `time` group, but nothing strange is happening in here. The number of crashes through time might be an interesting, as well as the number of crashes over holidays, but we would need to normalize it per day to know if the frequency is any greater.


The no holiday could use a better name than just empty string, so I recode it.

```
data[time] |> lapply(table, useNA="ifany")
```

```
## $crashYear
##
## 2000 2001 2002 2003 2004 2005 2006 2007 2008 2009 2010 2011 2012
## 31996 36125 38045 37950 37051 38364 39778 41661 39535 38247 36870 32450 30443
## 2013 2014 2015 2016 2017 2018 2019 2020 2021 2022 2023
## 30109 29784 32103 37249 39314 38469 36919 32808 34080 27982 4412
##
## $holiday
##
##      Christmas New Year      Easter      Labour Weekend
##      776922      20453      9463      7055
## Queens Birthday
```

```
##                                7851
data$holiday = replace(data$holiday, data$holiday == "", "Not holiday")

data$crashYear |> table() |>
  plot(las=2, ylab="", lwd=5, frame.plot=FALSE,
       main = "Number of crashes through time")
```



There is interesting cyclical behaviour. Obviously the 2023 year is incomplete, so we do not take that in account, and in similar fashion, I am not sure how complete are data from 2000, but otherwise there is a dip from 2009 to 2016. I don't know what happened at that time, since I arrived to NZ during 2014-2015 (so I single-handedly caused an increase of crashes, cool), and then again a decrease from 2019, which surely is due to Covid, but only on 2022 it we get global minimum, and 2022 is already after lockdowns. To be complete, 2007 is a global maximum of the number of crashes in our dataset.

Location variables Now to the location. There will be quite a few geographical variables, such as **region** and **subregion**, variables that describe the road where the crash happened, such as **crashLocation1** and **crashLocation2**, whether the crash happened in **urban** environment, and two direction variables, which I don't understand even after reading the key description, so I will drop them.

The **urban** variable is quite simple, it signify whether the crash happened in an urban environment. The crashes are about twice likely to happen in an urban environment compared to an open one. This is not that surprising as most people are concentrated in cities, but NZ is still relatively rural country and as we see later, most crashes happen on the State Highways, so the difference between urban and open crashes is not overwhelming. We will recode it from a categorical to binary variable.

```
data$crashDirectionDescription = NULL
data$directionRoleDescription = NULL
```



```
data$urban = setNames(c(0,1), c("Open", "Urban"))[data$urban]
```


Now to the **region** and **subregion**. In both cases, there are crashes that happened in the region marked as empty string. I assume that this means that the region is unknown and recode them as NA. But after looking at the data, I am not that sure about that, since their **crashLocation** is known. Again, expert knowledge would surely help and since the road and surely the dropped latitude and longitude are known, the region and subregion can be derived.

In total, there are 16 regions and 67 subregions. We can effectively visualize 16 regions, but 67 subregions would take some work.

```
data$region = with(data, {replace(region, region == "", NA)})
data$subregion = with(data, {replace(subregion, subregion == "", NA)})

source("src/graphics.r") # I prepared some customized barplots
data$region |> table() |>
  barplot(angle=35, cex=0.6, font=2, las=2,
    main="Crashes in across regions")
```

Crashes in across regions



```
data$subregion |> table() |> sort(decreasing=TRUE) |> head()
```

```
##
##      Auckland Christchurch City      Wellington City      Hamilton City
##      285346      53011      32876      28594
##      Dunedin City      Tauranga City
##      24707      18973
```

Auckland has the most crashes, with Christchurch second and Wellington third. A similar situation is happening with regions, although here the Waikato region is second.

However, this is to be expected, isn't it? Auckland will have the most crashes because Auckland is the biggest baddest city in NZ. We need to normalize this by population to get some unexpected insight.

For this, I downloaded some StatsNZ data for population, unfortunately it is only for 2019-2021, but anything will do. After all, we are not interested in trend and population will surely not change dramatically in the past 20 years.

The only potential problem is to match region names, I already did some cleaning in the population data, but seems that the NZ crash database does not like Maori spelling. For instance, there is Manawatū-Whanganui Region, but only Manawatu District.

```
source("src/population.r")
population = get_population_year(2021)
```

```
regions = data$region |> table() |> names()
regions[!regions %in% names(population)]
```

```
## character(0)
```

```
subregions = data$subregion |> table() |> names()
subregions[!subregions %in% names(population)]
```

```
## [1] "Manawatu District" "Taupo District"      "Whakatane District"
## [4] "Whangarei District"
```

We need to fix these regions.

```
# base R is not really good at this, so here is a good solution from my pkg:
replace2 = function(x, values, replace, ...){
  if(length(values) != length(replace))
    stop("The vector `values` and `replace` must have the same length!")

```

```
  match = match(x, values)
  x[!is.na(match)] = replace[match][!is.na(match)]
  x
}
```


```
before = c("Manawatu District", "Taupo District", "Whakatane District", "Whangarei District")
after = c("Manawatū District", "Taupō District", "Whakatāne District", "Whangārei District")
data$subregion = replace2(data$subregion, before, after)
```

Now we can normalize the regional counts per population:

```
regions = data$region |> table() |> names()
subregions = data$subregion |> table() |> names()
```

```
data$region |> table() |>
  (function(x){x/population[regions]}()) |>
  barplot(angle=35, cex=0.6, font=2, las=2, labels=FALSE,
    main = "Crashes across regions normalized by population")
```

Crashes across regions normalized by population



```
data$subregion |> table() |>
  (function(x){x/population[subregions]})( ) |>
  sort(decreasing=TRUE) |> head()
```

```
##
##   Waitomo District   Ruapehu District   Kaikoura District   Westland District
##      0.3231328      0.3065116      0.2948357      0.2676768
## Mackenzie District   Wairoa District
##      0.2594891      0.2522124
```

Now, that is interesting. Auckland is dethroned as the most dangerous city and instead West Coast Region is a hellhole! And Waitomo district, which I never heard about. Apparently, it sparsely populated rural area. It might be possible that there is a highway and as we will saw, most crashes happens on highway. So lets check it out.

```
WaitomoDistrict = subset(data, subregion == "Waitomo District")
WaitomoDistrict$crashLocation1 |> table() |> sort(decreasing=TRUE) |> head()
```

```
##
##      SH 3      SH 4      SH 30      SH 37 TE ANGA ROAD      RORA ST
##      1409      372      315      86      73      64
```

Looks like the hunch was confirmed. Most crashes happened on State Highways 3, 4 and 30.

Finally, we will explore the `crashLocation1` and `crashLocation2`. According to the key, the `crashLocation1` is the primary road where the crash happened, while the `crashLocation2` being the secondary one, side road, or place nearby. I don't know what exactly it means, but in both cases, State Highways and other long roads are the most frequent places for crashes.

Again, this should not be much of a surprise, since these roads are very long, so assuming that the chance that the car crashes is the same for every single road, the longer the road, the more crashes we would expect to see.

Unfortunately, there is not a good way for me to check this assumption and normalize the number in the same way as we did it for the regions with population. I doubt that there is a statistics that specifies how long each road is.

```
data$crashLocation1 |> table() |> length() # 37453 unique values
```

```
## [1] 37453
```

```
data$crashLocation2 |> table() |> length() # 52060 unique values
```

```
## [1] 52060
```

```
data$crashLocation1 |> table() |> sort(decreasing=TRUE) |> head()
```

```
##
##          SH 1N          SH 2          SH 1S          SH 3
##          54779          21225          17586          10898
## GREAT SOUTH ROAD          SH 6
##          10243          9834
```

```
data$crashLocation2 |> table() |> sort(decreasing=TRUE) |> head()
```

```
##
##          SH 1N GREAT SOUTH ROAD          SH 2 GREAT NORTH ROAD
##          3765          3340          2337          1822
##          QUEEN ST          SH 1S
##          1771          1742
```

Weather variables The three weather variables are `light`, `weatherA` and `weatherB`.

I am not sure if the `Dark` category in the `light` variable means nighttime, or if this might mean that it was just very cloudy day, but I guess cloudy day is the `Overcast` category. In that case, about 30% of crashes happened during night. Since there is less traffic during night, this is quite significant increase from what I would expect, but I have no means of normalizing this to get precise answer.

```
data$light |> table(useNA="ifany") |> (function(x){x/sum(x)})()
```

```
##
## Bright sun          Dark          Overcast          Twilight          Unknown
## 0.368695117 0.274927471 0.299935260 0.046695808 0.009746344
```

```
data$light[data$light == "Unknown"] = NA
```

```
data$weatherA |> table(useNA="ifany")
```

```
##
##          Fine Hail or Sleet          Heavy rain          Light rain          Mist or Fog
##          635621          132          33153          124210          11306
##          Null          Snow
##          15778          1544
```

```
data$weatherA[data$weatherA == "Null"] = NA
```

The two other weather variables are quite peculiar. `weatherA` encodes weather, everything seems to be relatively standard, just using `Null` to encode unknown data (that is like the fifth different value for unknown data we have encountered, bruh). The majority of crashes happened during `Fine` weather and `Light rain`,

which would be the majority of days, it rains quite a lot in NZ, while Hail and Snow are quite uncommon. Still, there was quite a lot of fog in Dunedin so I would expect more crashes during that time. We could check it out:

```
Foggy = subset(data, weatherA == "Mist or Fog")
Foggy$region |> table() |> sort(decreasing=TRUE)
```

```
##
##           Waikato Region           Auckland Region           Canterbury Region
##                2566                2154                1699
##           Otago Region           Northland Region           Wellington Region
##                932                615                613
## Manawatū-Whanganui Region       Bay of Plenty Region           Southland Region
##                506                489                447
##           Hawke's Bay Region       West Coast Region           Taranaki Region
##                444                224                196
##           Tasman Region           Gisborne Region           Marlborough Region
##                127                112                108
##           Nelson Region
##                52
```

```
Foggy$subregion |> table() |> sort(decreasing=TRUE) |> head()
```

```
##
##           Auckland Christchurch City Waikato District           Hamilton City
##                2154                800                645                484
##           Dunedin City Far North District
##                458                293
```

Looks like Waikato is the most misty region, at least regarding crashes. Auckland is still there, although it is not that dominant as we would expect from the total number of crashes and population. All things considered, Dunedin is not that misty, even Christchurch beats it. But again, if we normalized this by the total number of crashes or population, the numbers would surely look different. This doesn't mean that this pattern doesn't exist and it is not a valuable insight, just that there are different ways how to look on these patterns.

Now, **WeatherB** is quite a bit weirder. It has two categories **None** and **Null**. Not sure if it means unknown data or just nothing further description. The overwhelming presence of the **Null** would suggest just no further information. Given this, I am not that willing to use this variable for further modelling. Frost and Strong Wind would certainly have an effect on the probability that a crash will occur or on their severity, but this variable is miscoded. Merging it with **WeatherA** would be meaningful, although it would introduce a bit too many categories, which is another reason for dropping this, or make a note to look at this later with an expert.

```
data$weatherB |> table()
```

```
##
##           Frost           None           Null Strong wind
##           9254             5       798096       14389
```

```
lapply(c("Frost", "Strong wind"), function(x){subset(data, weatherB == x, select=weatherA) |> table()})
```

```
## [[1]]
##
##           Fine Hail or Sleet           Heavy rain           Light rain           Mist or Fog
##           7140                21                43                333                894
##           Snow
##           385
```

```
##
## [[2]]
##
##      Fine Hail or Sleet      Heavy rain      Light rain      Mist or Fog
##      7217              23              3274              3360              135
##      Snow
##      177
```

When we look at interactions between `weatherA` and `weatherB`, we can see that `Frost` is associated a little bit more with `Mist or Fog` and `Snow`, while `Strong wind` is associated quite strongly with `Heavy rain` and `Light rain`. In both cases, the number of crashes during `Fine` is almost identical.

Crash Severity variables Now we are getting into interesting data that we might explore rather with modelling.

```
lapply(data[crashSeverity], table, useNA="ifany")
```

```
## $crashSeverity
##
##      Fatal Crash      Minor Crash Non-Injury Crash      Serious Crash
##      7589              191336              575954              46865
##
## $fatalCount
##
##      0      1      2      3      4      5      6      7      8      9      <NA>
## 814154  6854  567  115  39   7   3   2   1   1   1
##
## $minorInjuryCount
##
##      0      1      2      3      4      5      6      7      8      9      10
## 615625 165582 30358  6996 2164  649  228  83  23  13   7
##      11      12      13      14      15      16      18      21      26      30      34
##      2      1      1      1      2      2      2      1      1      1      1
##      <NA>
##      1
##
## $seriousInjuryCount
##
##      0      1      2      3      4      5      6      7      8      9      10
## 772759 43060 4596  924  271  86  28  8  5  1  3
##      12      14      <NA>
##      1      1      1
```

The `crashSeverity` variable is a summary variable that tells us how severe was the crash. Fortunately, only 7589 crashes were fatal over the 20 years, with total of 8573 people perished. The most serious was a crash where 9 people died in total.

```
# Total number of people died over the 20 year period:
```

```
data$fatalCount |> table() |> (function(x){x * as.numeric(names(x))})() |> sum()
```

```
## [1] 8573
```

```
# Explore the cases where deaths > 6
```

```
# subset(data, crashSeverity > 6)
```

Looking at the 4 cases with high death count, I don't see anything particular. In all cases, this happened on fine weather, State Highway with sealed road with speed limit 100. In one case there was temporary speed

limit 30, maybe due to roadwork but the `roadwork` variable is set to unknown.

Other objects Category of “I don’t know where to put it”.

```
lapply(data[otherObjects], table, useNA="ifany")
```

```
## $debris
##
##      0      1      2      3      4      5      6      7  <NA>
## 330374  2376   131   24     4     2     1     1 488831
##
## $objectThrownOrDropped
##
##      0      1      2      3      4  <NA>
## 332232   636    35     8     2 488831
##
## $otherObject
##
##      0      1      2      3      4      5  <NA>
## 325189   7662    51     7     2     2 488831
##
## $pedestrian
##
##      1      2      3      4      5      6  <NA>
## 25681    785   110    23     3     3 795139
##
## $phoneBoxEtc
##
##      0      1      2      3  <NA>
## 328797   4096    19     1 488831
##
## $postOrPole
##
##      0      1      2      3      4  <NA>
## 292252   40439   214     7     1 488831
##
## $strayAnimal
##
##      0      1      2      3  <NA>
## 331843    994    72     4 488831
##
## $tree
##
##      0      1      2      3  <NA>
## 299466   33089   354     4 488831
```

From the pattern of NA values, it looks like that this category is quite interconnected, with only `pedestrian` being the weird one. The `pedestrian` variable should probably be in the `vehicles` class.

Out of all of these, the only common objects seem to be the `postOrPole` and `tree` with about 5 and 4 percent of all crashes respectively. Other objects are quite rare.

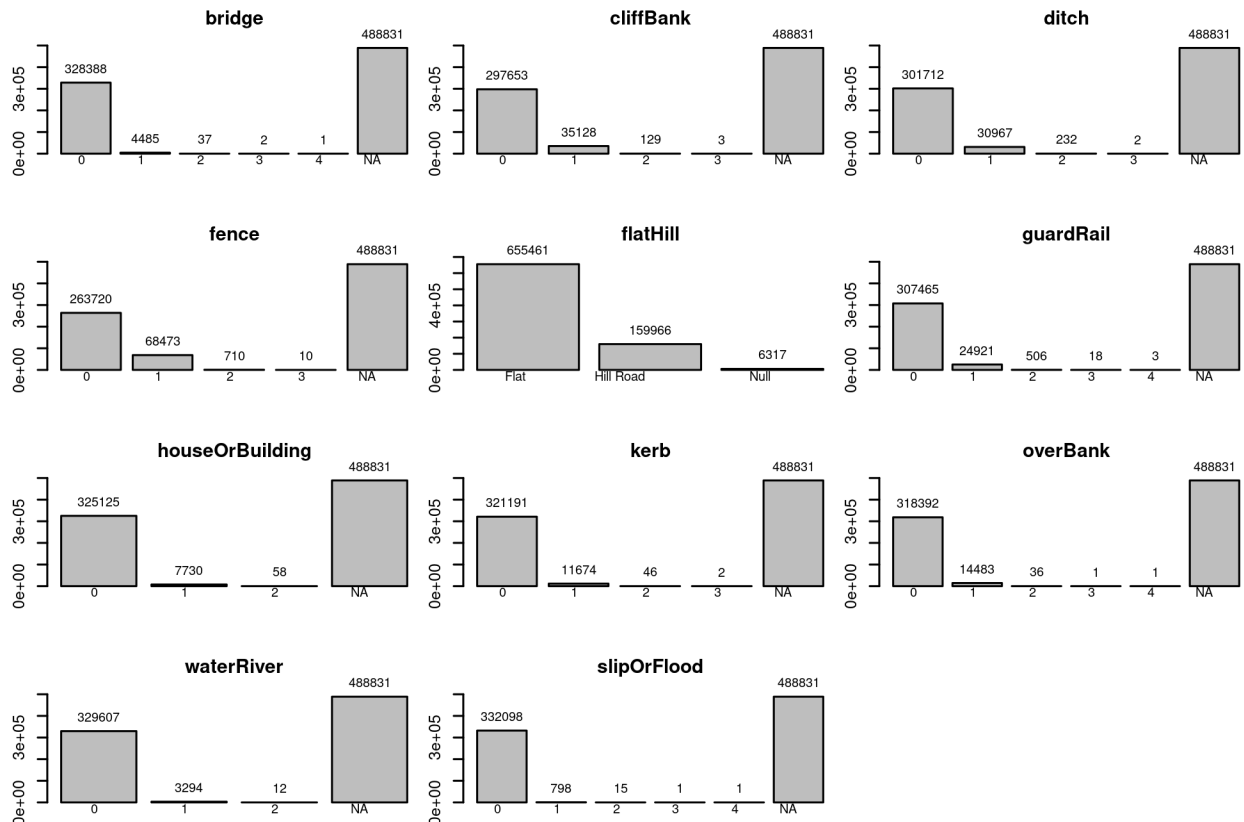
Terrain features First thing that hits me is the pattern of NA. We already saw it in the `otherObjects` category. Looks like for more than half of the data, any further description is simply missing. So I looked again at the key and these all are not description of the state of the terrain, but *how many times X was hit*

during the crash. That is, the fence was hit once in 68473 cases. Other than fence, only `CliffBank` and `ditch` appear to be somewhat common, but well below 5 percent of cases.

The only description of terrain is in fact `flatHill`, which is also a badly encoded variable with categories `Flat`, `Hill Road` and `Null`. So I make null about all this, remove the `flatHill` and encode it simply as a hill with 0, 1 and NA cases.

```
# Its quite bit of text, so we won't use this:
# lapply(data[terrainFeature], table, useNA="ifany")
```

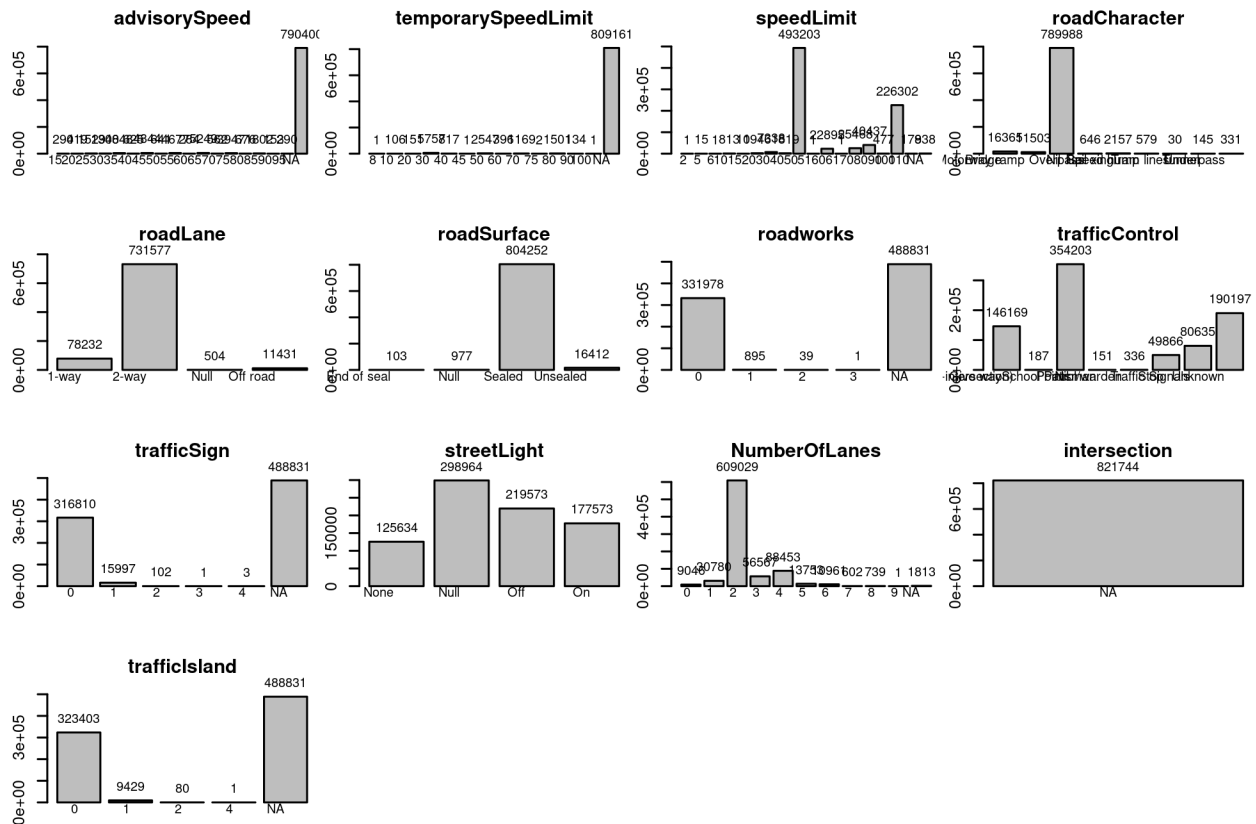
```
# Instead we use barplots from src/graphics.r
data[terrainFeature ] |> lapply(table, useNA="ifany") |> barplots()
```



```
data$hill = setNames(c(0, 1, NA), c("Flat", "Hill Road", "Null"))[data$flatHill]
data$flatHill = NULL
```

Road conditions Remember, this is EDA, the plots do not have to be pretty or even that detailed. When we are looking cursory on a large number of plots, all we need to know is if there is some basic pattern and if we need to look deeper.

```
data[roadConditions ] |> lapply(table, useNA="ifany") |> barplots()
```

For instance, we can already see that `intersection` is degenerated/constant variable and can be simply dropped. We can also see that we need to review Speed variables (`advisorySpeed`, `temporarySpeedLimit`, and `speedLimit`) separately, look at `roadCharacter` and `trafficControl` more closely, and recode a bunch of Null into NA.

So let's look at the variables more closely.

```
data[c("advisorySpeed", "temporarySpeedLimit", "speedLimit")] |>
  lapply(table, useNA="ifany")
```

```
## $advisorySpeed
##
##      15      20      25      30      35      40      45      50      55      60      65
##    290    419    1529    1308    4048    625    4844    641    4677    284    5249
##      70      75      80      85      90      95    <NA>
##    562    3947    676    1802    153    290 790400
##
## $temporarySpeedLimit
##
##       8      10      20      30      40      45      50      60      70      75      80
##       1     106     151    5757     817      1    2547    396    1169      2    1501
##      90     100    <NA>
##     134       1 809161
##
## $speedLimit
##
##       2       5       6      10      15      20      30      40      50      51      60
```

```
##      1      15      1    813     10   1946   7338   1819 493203      1 22895
##     61     70     80     90    100    110   <NA>
##      1 25468 40437  477 226302   179   838
```

```
data$advisorySpeed = NULL
data$temporarySpeedLimit = NULL
data$intersection = NULL
```

The number of NA values in `advisorySpeed` and `temporarySpeedLimit` suggest that the data are missing not because they are unknown, but because there is no advisory or temporary speed. This is another point we need to be aware when doing modelling, as interpretation will change drastically.

For instance, since this value is not missing, it would be folly to use methods to estimate it, such as through frequentist or bayesian mean. Instead, it is a conditional variable, either there is or isn't advisory/temporary speed, or there is one with a certain value.


I got a feeling that trying to model these variables in relationship to `speedLimit` would be something that would most resemble the conditions on road, i.e., looking at what is the speed limit and if it is significantly different from the other two variables. If it is, flipping a switch in a binary representation. For instance, if there is a sharp turn with advisory speed 25, it will mean something different on State Highway with speed limit 100 and smaller road with speed limit 30.

But this is really something that should be discussed with an expert. After all, this is why we are doing EDA after all. In the meantime, I will not use these variables (i.e., drop them), since properly exploring them will take quite a lot of time, I am already spending too much time on this, and I still can't see the end.

Now back to the `speedLimit`. I did not honestly know that there were places with speed limit 2. Looking at the item with `subset(data, speedLimit == 2)` tells me that it is a crash in Auckland during nighttime on a sealed road between two station wagons, one of which was parked, but no temporary speed limit or roadwork.

Next we look at `roadworks`, `trafficSign` and `trafficIsland`. Unlike what I originally thought, these are items similar to a `tree`, `riverBanks` and similar, objects that were hit during the crash, and not road conditions. Again, not very interesting variables since we lack this information for more than half of the data points. I will again make note for a future me here and write down some thoughts. Feels to me like we could derive a single variable from all of these, such as "other objects were hit". But it is also likely that these variables were collected for a reason, I can imagine crashes against trees being absolutely lethal, and encoding trees together with other variables might not be what we want. So maybe binarize them all and use Lasso to find which ones are significant?

```
data[c("roadworks", "trafficSign", "trafficIsland")] |> lapply(table, useNA="ifany") |> barplots()
```



Road character is another variable that could use better encoding. There is a big category Nil cotaining the majority of the data. But it looks like No special feature or Normal rather than Unknown data. In here the encoding is quite obvious so I will recode it this way.

```
data$roadCharacter |> table(useNA="ifany")
```

```
##
##      Bridge Motorway ramp      Nil      Overpass      Rail xing
##      16365      11503      789988      646      2157
##      Speed hump      Tram lines      Tunnel      Underpass
##      579      30      145      331
```

```
data$roadCharacter = replace2(data$roadCharacter, "Nil", "Normal")
```

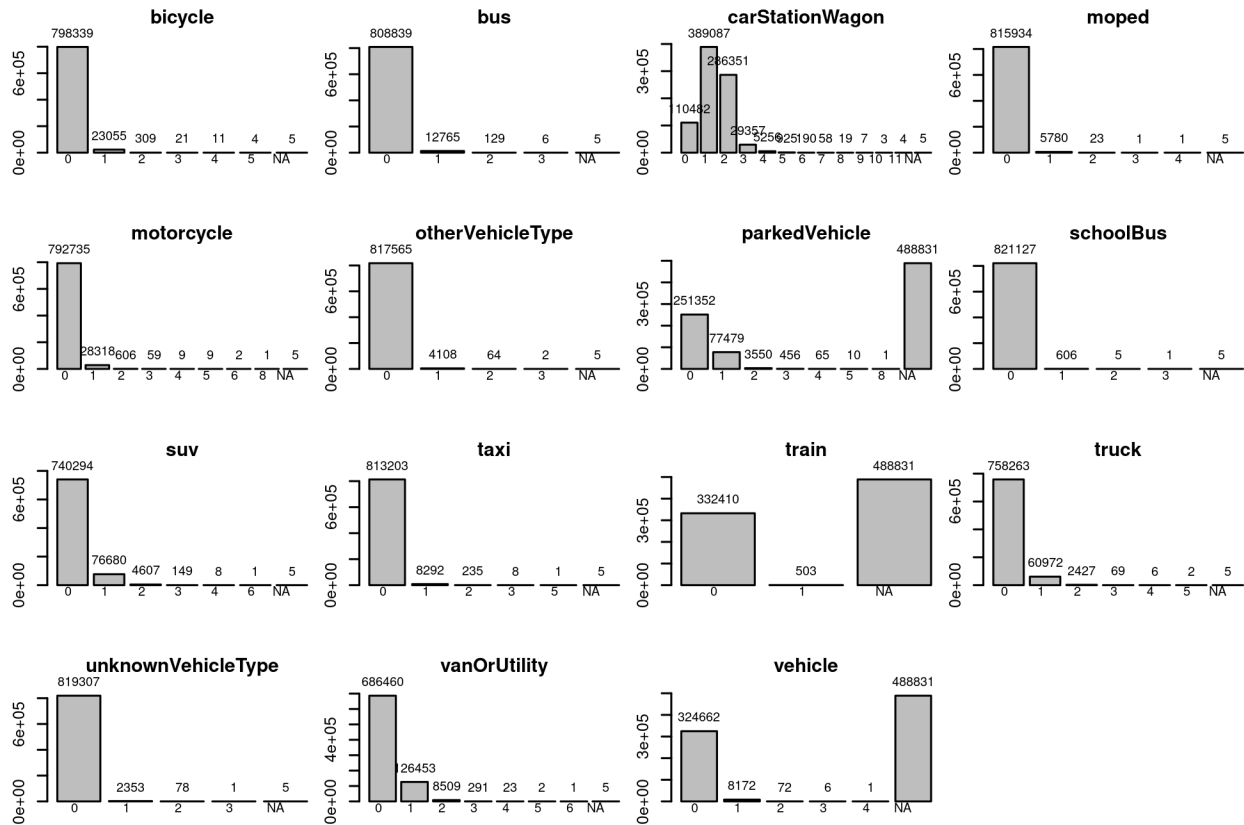
trafficControl has not one, but two different “missing data” variables, Nil and Unknown. The key is not really helpful as it just says that these categories exists, but what is the difference between those two? For safety, I am dropping this variable.

We are left with roadLane, roadSurface, streetLight, and numberOfLanes. They looks well-behaved, we just need to recode missing data for three of them.

```
data$trafficControl = NULL
data$roadLane = replace2(data$roadLane, "Null", NA)
data$roadSurface = replace2(data$roadSurface, "Null", NA)
data$streetLight = replace2(data$streetLight, "Null", NA)
```

Vehicle Type And we are at the last variable class, the vehicleType class.


```
data[vehicleType] |> lapply(table, useNA="ifany") |> barplots()
```



From the missing values, you can see that there are two classes. The `parkedVehicle`, `train` and `vehicle`, and the rest. The first tree are interpreted as “how many times X was struck during the crash”, while the other variables are keyed as “how many X were involved in the crash”. Quite the different interpretation, and the pattern of missing variables. I can feel something is happening in there, we saw the 488831 missing data quite a lot.

```
# Not super readable, but you can see the general pattern
is.na(data) |>
  colSums() |>
  sort(decreasing=TRUE) |>
  barplot(angle=45, labels=FALSE, cex.axis=0.4, font=2,
    main = "Missing data across variables")
```

Missing data across variables



For the list of variables:

```
is.na(data) |> colSums() |> (function(x){subset(x, x == "488831")})() |> names()
```

```
## [1] "bridge"           "cliffBank"        "debris"
## [4] "ditch"            "fence"            "guardRail"
## [7] "houseOrBuilding"  "kerb"             "objectThrownOrDropped"
## [10] "otherObject"      "overBank"         "parkedVehicle"
## [13] "phoneBoxEtc"      "postOrPole"       "roadworks"
## [16] "slipOrFlood"      "strayAnimal"      "trafficIsland"
## [19] "trafficSign"      "train"            "tree"
## [22] "vehicle"          "waterRiver"
```

All of them are other objects that might have been hit during the crash.

And this is all. We went through all the variables and looked at their individual patterns. Now we should move to relationship between two variables, but I run out of time, so maybe later!

Patterns across regions


We do not have time to do full bivariate exploration, but we wanted to look at some patterns across regions.

First, I want to look at which region is the deadliest. Don't worry, I don't have any unhealthy morbid obsession, but since we are investigating car crashes, this seems a natural thing to look at and try to identify causes. We already know that Auckland has the most crashes simply because it has the highest population, so we will normalize according to the number of crashes.

```
crashSeverity = table(data$region, data$crashSeverity)[
  , c("Non-Injury Crash", "Minor Crash", "Serious Crash", "Fatal Crash")
] # sorted for convenience
```

```
col = palette.colors(4, "Set 1")


par("mar" = c(6,9,2,2))
(crashSeverity / rowSums(crashSeverity)) |>
  t() |>
  graphics::barplot(horiz=TRUE, las=1, col=col, border=NA, cex.names=0.6, font=2,
    main = "Severity of crashes")
legend("top", legend = colnames(crashSeverity), fill = col, ncol = 2,
  inset = c(0, 1.1), xpd = TRUE, bty = "n")
```



Contrary to the number of crashes, Auckland seems to be a relatively safer region, since most crashes are without any injury. On the other hand, West Coast and Northland are more dangerous

```
par("mar" = c(4,9,2,2))
fatalCount = table(data$region, data$fatalCount)
(t(fatalCount) * as.numeric(colnames(fatalCount))) |>
  colSums() |>
  graphics::barplot(horiz=TRUE, las=1, border=NA, cex.names=0.6, col=col[4],
    font=2, main="Total crash fatalities")
```

Total crash fatalities

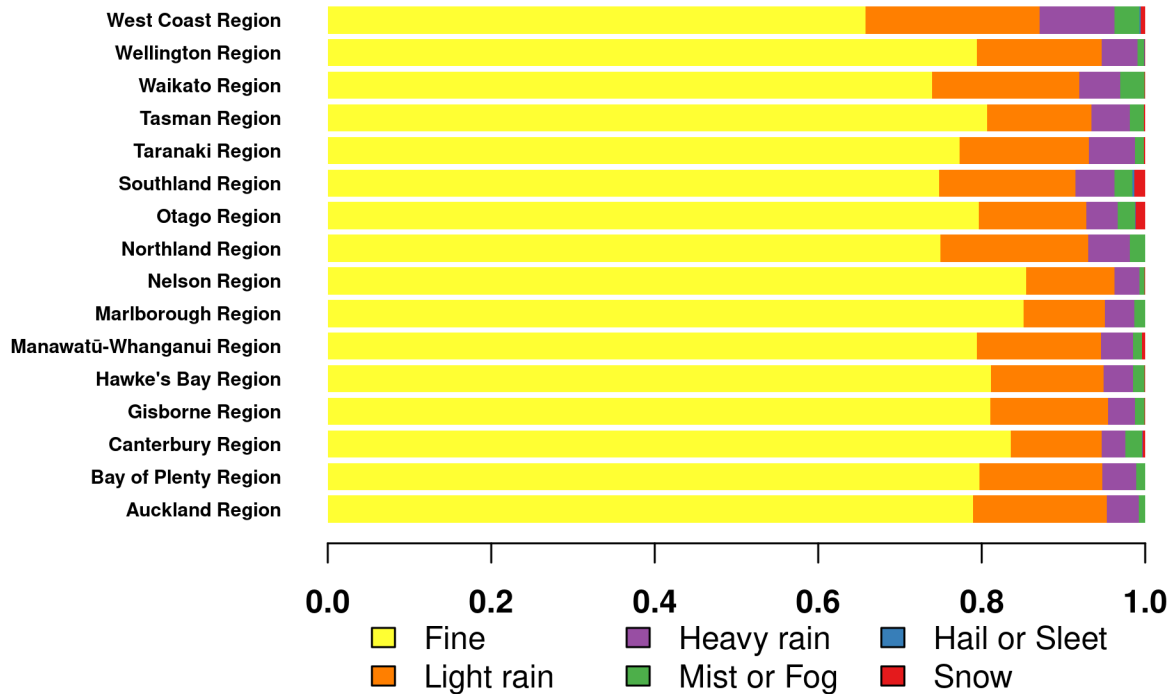


Another way to look at this is to look at total number of deaths, this include multiple deaths per crash. Despite Waikato having only 30% population of Auckland, it is way above Auckland in the number of deaths in crashes over 20 year period.

And to finish it, to escape the morbidity, I will look at weather.

```
weather = table(data$region, data$weatherA)[
  , c("Fine", "Light rain", "Heavy rain", "Mist or Fog", "Hail or Sleet", "Snow")
] # sorted for convenience
col = palette.colors(6, "Set 1") |> rev()
par("mar" = c(6,9,2,2))
(weather / rowSums(weather)) |>
  t() |>
  graphics::barplot(horiz=TRUE, las=1, col=col, border=NA, cex.names=0.6, font=2,
    main = "Weather during accidents")
legend("top", legend = colnames(weather), fill = col, ncol = 3,
  inset = c(0, 1.1), xpd = TRUE, bty = "n")
```

Weather during accidents



Looks like Nelson and Marlborough are the sunniest regions, while Southland and Otago the snowiest. I can confirm, I lived in Dunedin, on the hills there. But usually the snow thawed during day, so you could enjoy it only at night or early in the morning. Still, snow in New Zealand, fun times.

Modelling

Here we just do a quick look into the modelling. Some simple models are convenient to do during data-exploration, particularly the three methods are very convenient, be it CART or RandomForest. They can all handle continuous and categorical data, missing data, and are relatively robust, with RandomForest being also quite performant, often just little behind well-optimized Gradient Boosting algos.

A big advantage of CART is that it provides a great graphical output and is easy to interpret. Although they are not as performant, they can help you discover patterns in data, which means they are great for EDA.

First of all, we will make a model data containing a variables we want to use during modelling.

Then we fit a simple CART. At first, we do not care about any test error, we want to just ascertain the pattern in the data.

We want to model `crashSeverity` based on various reasonable variables we have selected during our process. Lets start with types of vehicles involved, weather, light conditions, holiday time, hill, urban environment, road characters, and whether, but not other objects.

We won't include `fatalCount`, `minorInjuryCount` and `majorInjuryCount` as this will cause data leakage.

```
mdata = data[c(
    "bicycle", "bus", "carStationWagon",
    "crashSeverity", "crashSHDescription",
    "holiday", "light", "moped", "motorcycle",
    "NumberOfLanes", "otherVehicleType", "pedestrian",
```




```

    "roadCharacter", "roadLane", "roadSurface",
    "schoolBus", "speedLimit", "streetLight",
    "suv", "taxi", "truck", "unknownVehicleType",
    "weatherA", "hill"
  )]
# we also need to convert them to factors
chars = lapply(mdata, class) == "character"
mdata[chars] = lapply(mdata[chars], factor)

library("tree")
tree = tree(crashSeverity ~ ., data=mdata)
plot(tree)
text(tree, pretty=0, cex=0.5)

```



Uh, I have expected a more deeper tree. This is really bad, even if the model is ultimately not good, you would expect for CART to pick up some pattern.

The error is terrible, we have predicted over 80% of classes wrong.

We can see what is happening by looking at the contingency table. Remember that because of our mad model, we are classifying only between minor and serious crash.

```

predicted = predict(tree, newdata = mdata, type="class")
error = predicted != mdata$crashSeverity
table("Predicted"=predicted, "actual"=data$crashSeverity)

```

```

##          actual
## Predicted  Fatal Crash Minor Crash Non-Injury Crash Serious Crash

```

## Fatal Crash	0	0	0	0
## Minor Crash	1797	113086	391522	22476
## Non-Injury Crash	0	0	0	0
## Serious Crash	5792	78250	184432	24389

Looks like we just can't predict well non-injury crashes. These are distributed both among our predicted minor and serious crash.

We know that CART is bad, but we didn't know that it is that bad. We would expect it to pick up on some signal, but either we have removed it by filtering some features, or type of crash just can't be predicted from the data on hand.

But here is a thought. Non-injury is quite overrepresented in the data and thus we would need to build quite deep tree to find differences. I can get to this point my manipulating with `mindev`, so maybe there is something hidden deep inside, but there isn't a single clear pattern.

```
library("randomForest")
```

```
## randomForest 4.7-1.1
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
set.seed(1) # for replicability
```

```
mdata = na.roughfix(mdata) # randomForest doesn't handle missing values well
```

```
mdata = mdata[sample(1:nrow(mdata), size=20000),] # subsample, otherwise my computer breaks
```

```
rf = randomForest(crashSeverity ~ ., data=mdata)
```

```
rf
```

```
##
```

```
## Call:
```

```
## randomForest(formula = crashSeverity ~ ., data = mdata)
```

```
##           Type of random forest: classification
```

```
##           Number of trees: 500
```

```
## No. of variables tried at each split: 4
```

```
##
```

```
##           OOB estimate of  error rate: 27.78%
```

```
## Confusion matrix:
```

```
##           Fatal Crash Minor Crash Non-Injury Crash Serious Crash
```

```
## Fatal Crash           0           17           151           13
```

```
## Minor Crash           0           628           3861           76
```

```
## Non-Injury Crash      0           315          13739           22
```

```
## Serious Crash         1           254           846           77
```

```
##           class.error
```

```
## Fatal Crash          1.00000000
```

```
## Minor Crash          0.86243154
```

```
## Non-Injury Crash     0.02394146
```

```
## Serious Crash        0.93463497
```

The OOB estimate of error is promising, but closer look at the confusion matrix and class error shows that we are unable to predict anything and the relatively low error is caused purely by the overrepresented Non-Injury Crash.

Conclusion

I have performed Exploratory Data Analysis on the New Zealand Crash data. It was a bit contrived and we have spend a lot of time trying to clean the data and figure out what each variable or category means rather

than looking more into patterns. This is another reason why you should always have an expert on hand, or become one, when you are working with data.

Aside of cleaning, we have did some exploration, notably into the fatalities. Auckland is not as terrifying as it might look like, it has a large number of crashes, but this is to be expected due to its large population, and most crashes are without injury. On the other hand, you should be a bit worried if you are living or driving in Waikato.

In the end, we weren't able to find reason for why crashes become fatal. This is another reason why you should be vigilant and try to err on the safe side. Any crash, no matter the speed, road type, or weather condition can turn into serious injury or even become fatal. So drive safe, you are not the only one on the road.